

Intro to Crypto and Cryptocurrencies

Slides by Arvind Narayanan et al.


Hash Pointers and Data Structures

hash pointer is:

- * pointer to where some info is stored,
- and
- * (cryptographic) hash of the info

if we have a hash pointer, we can


- * ask to get the info back, and
- * verify that it hasn't changed



key idea:

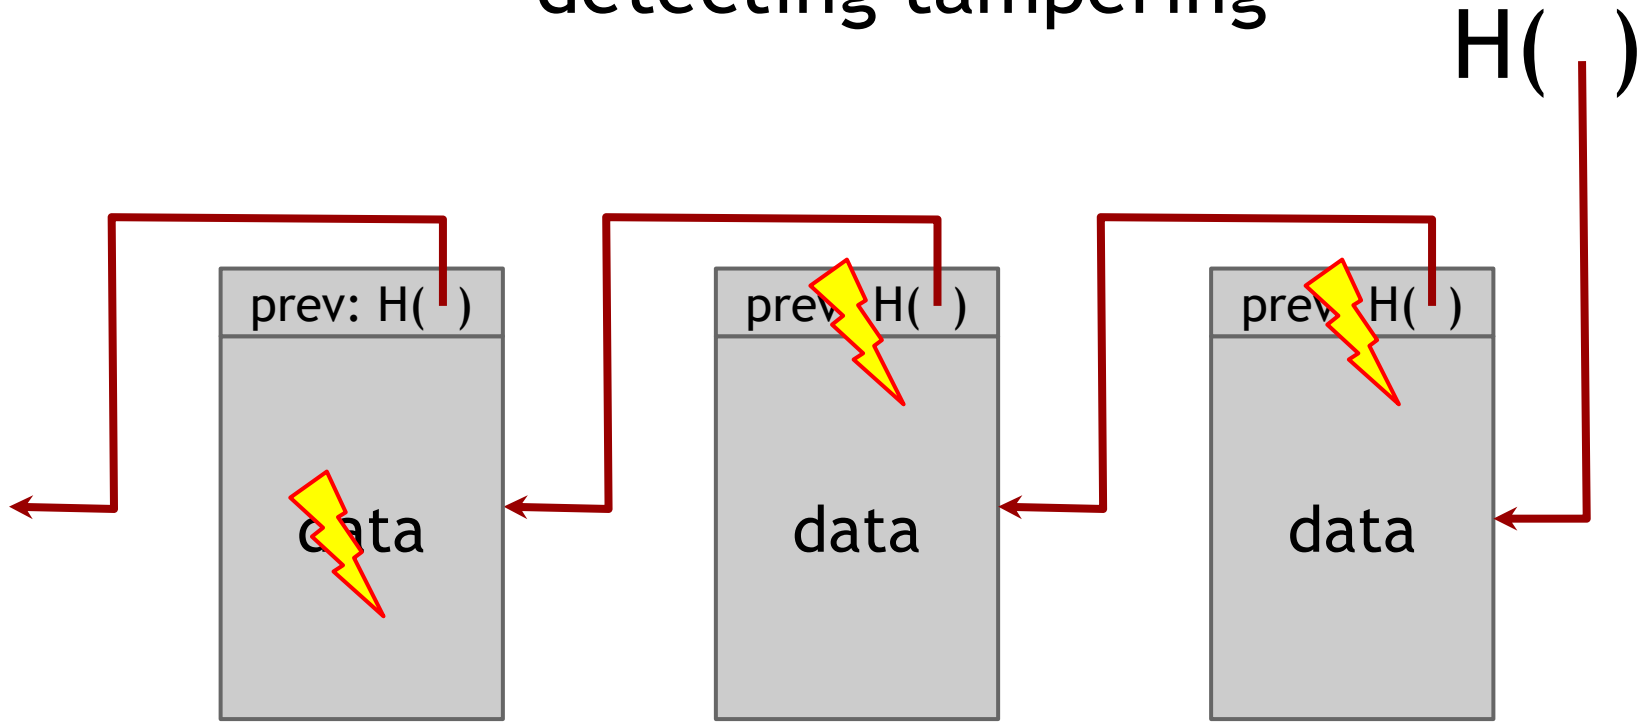
build data structures with hash pointers

linked list with hash pointers = “block chain”



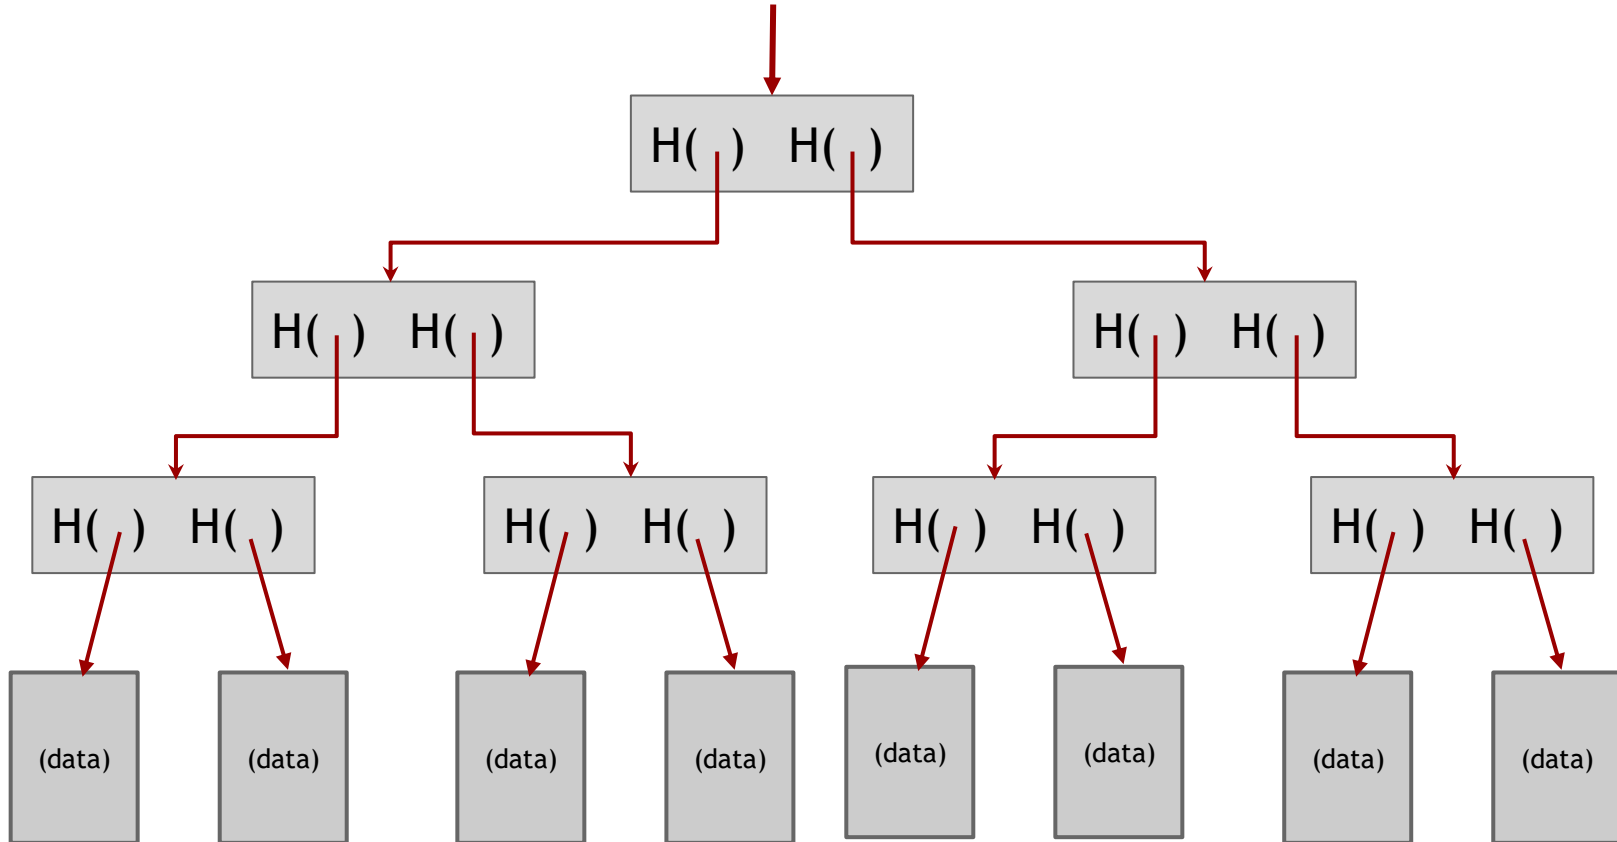
use case: tamper-evident log

detecting tampering

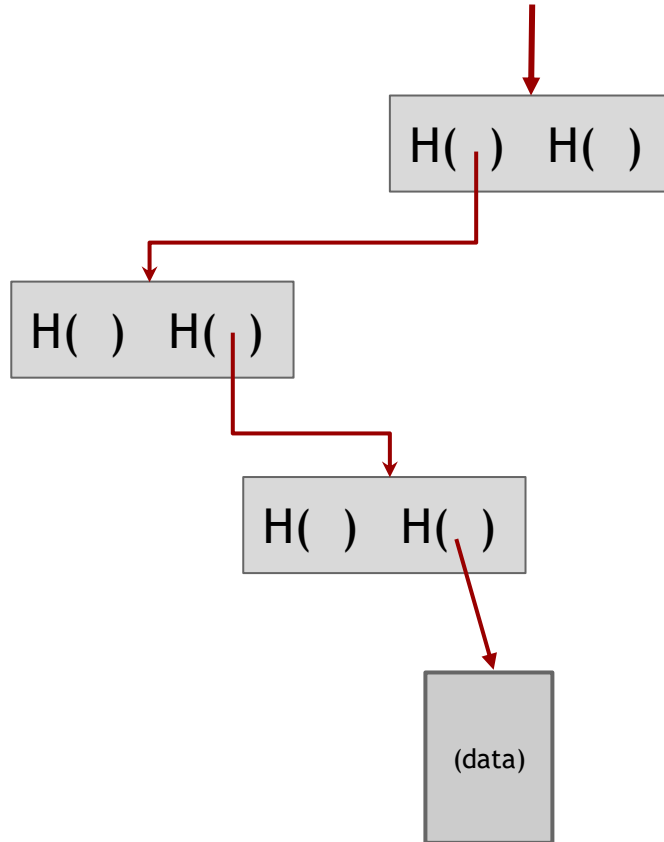


use case: tamper-evident log

binary tree with hash pointers = “Merkle tree”



proving membership in a Merkle tree



show $O(\log n)$ items

Advantages of Merkle trees

Tree holds many items

but just need to remember the root hash

Can verify membership in $O(\log n)$ time/space

Variant: sorted Merkle tree

can verify non-membership in $O(\log n)$

(show items before, after the missing one)

More generally ...

can use hash pointers in any pointer-based data structure that has no cycles



GoofyCoin

Simple Cryptocurrencies

Obvious approach

1. Use public keys as addresses
2. Sign to authorize transfer to new address

New coins created [somehow]

Goofy can create new coins

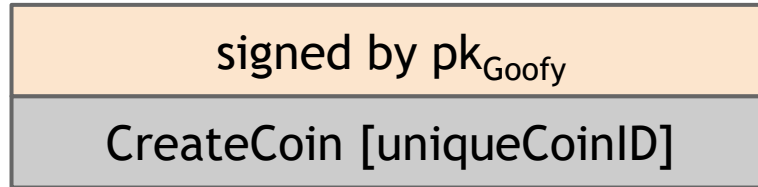
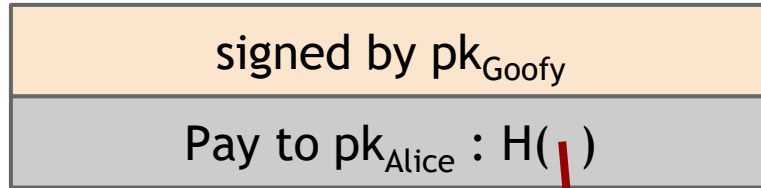
signed by pk_{Goofy}

CreateCoin [uniqueCoinID]

New coins belong to me.



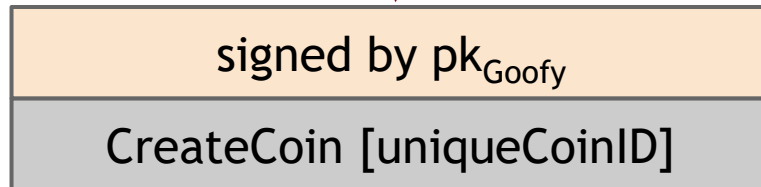
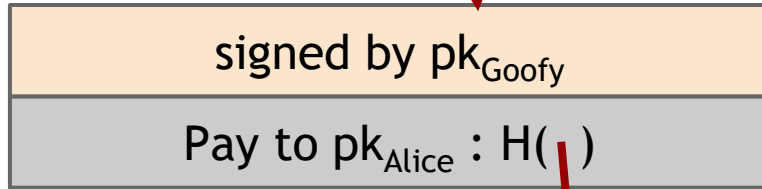
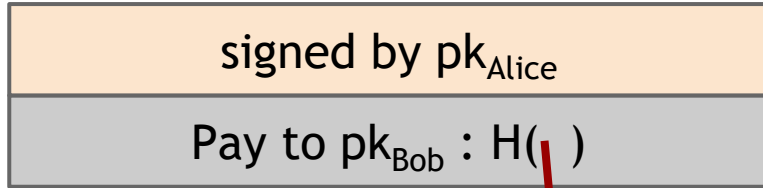
A coin's owner can spend it.



Alice owns it now.



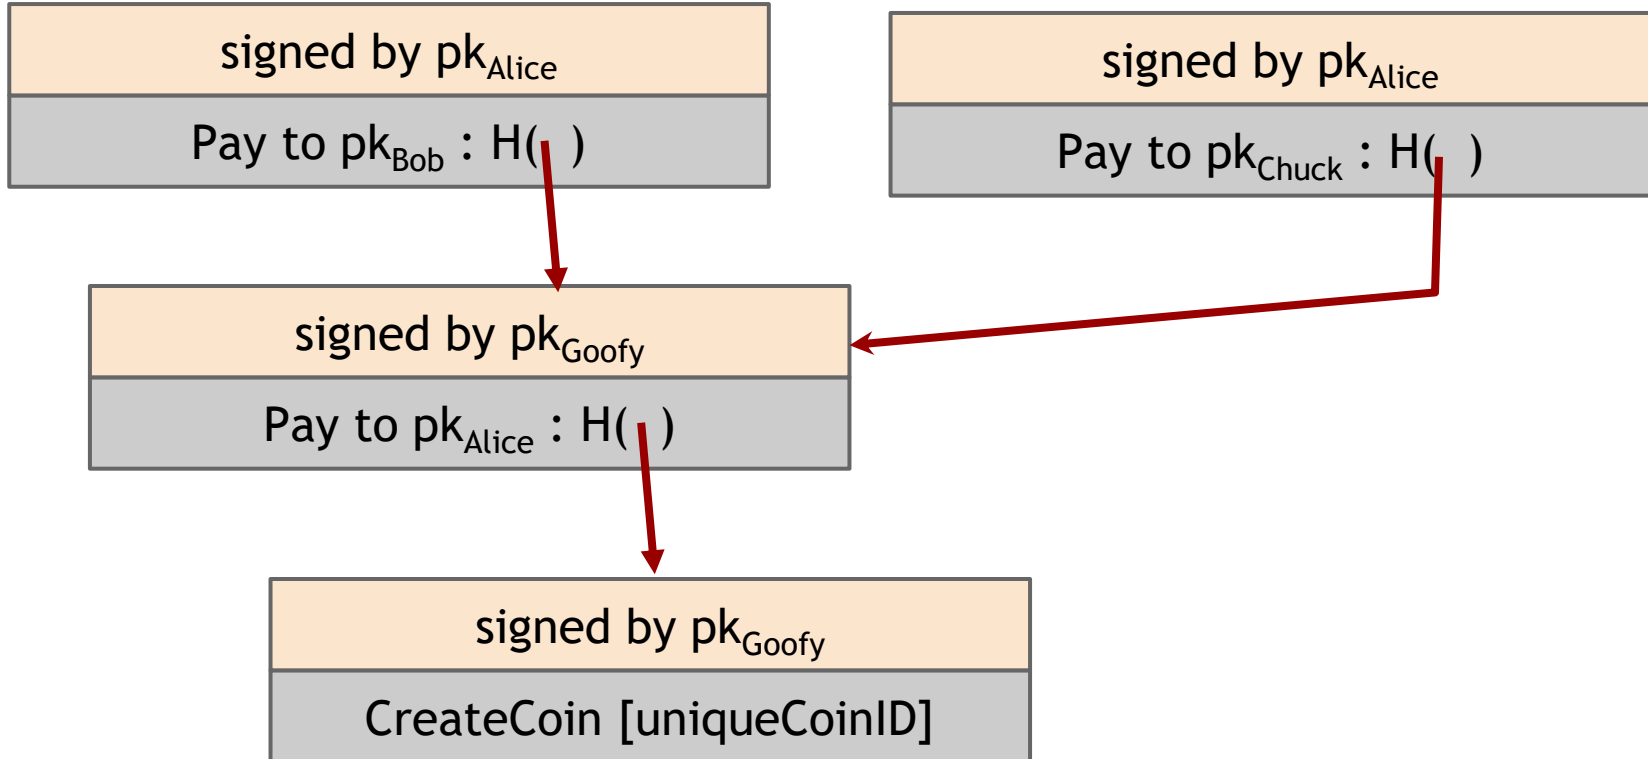
The recipient can pass on the coin again.



Bob owns it now.




double-spending attack



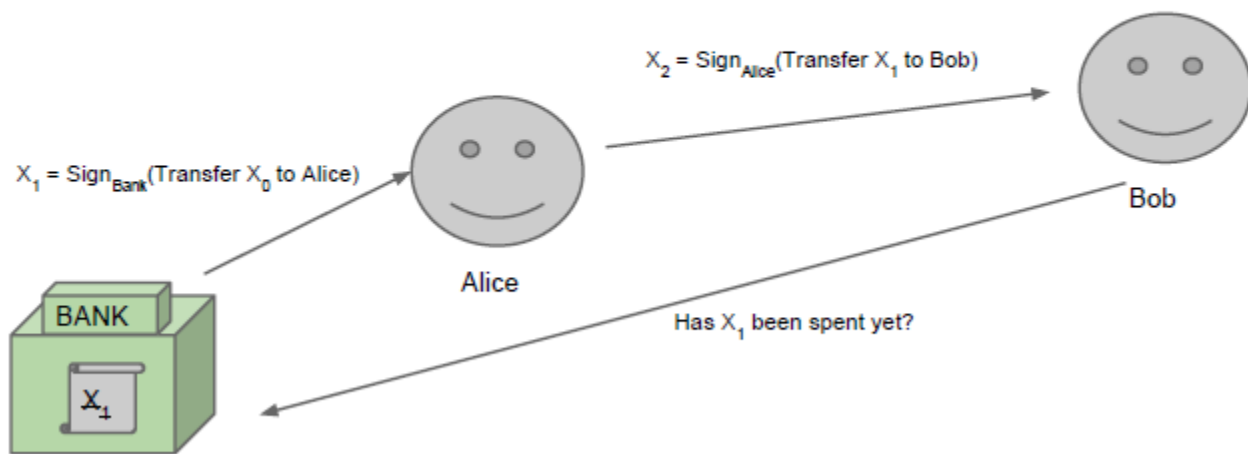
double-spending attack

the main design challenge in digital currency

Double-spends must be prevented



Traditional approach: talk to the issuer



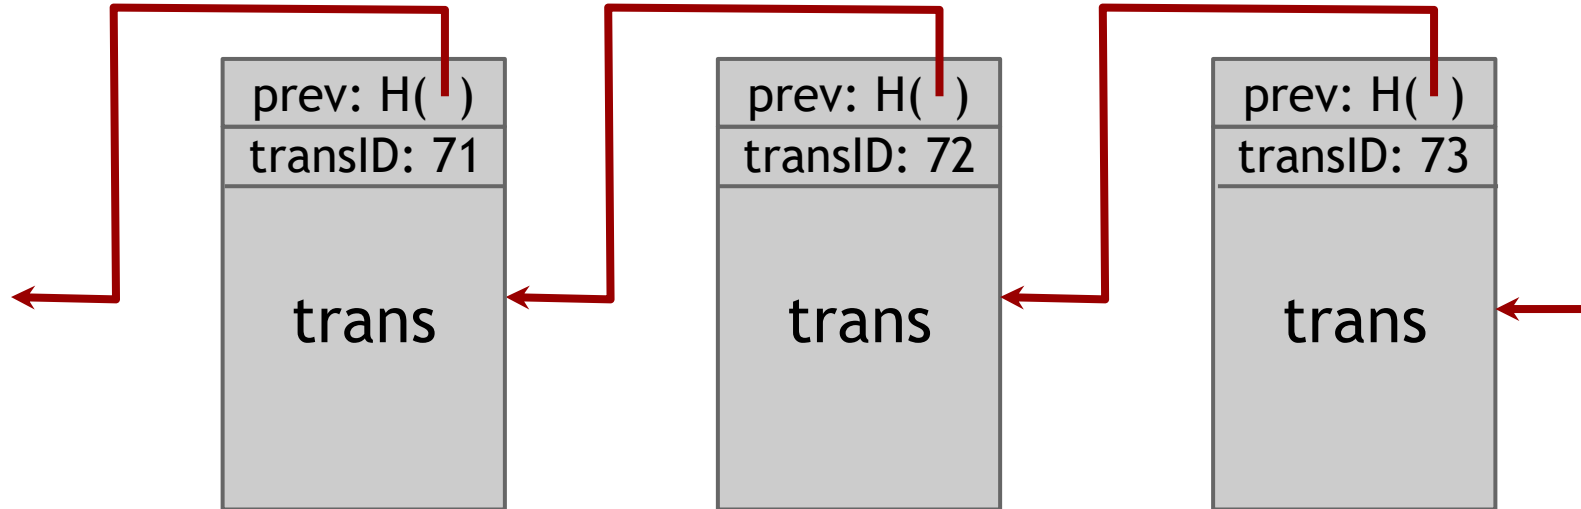


ScroogeCoin

Scrooge publishes a history of all transactions
(a block chain, signed by Scrooge)

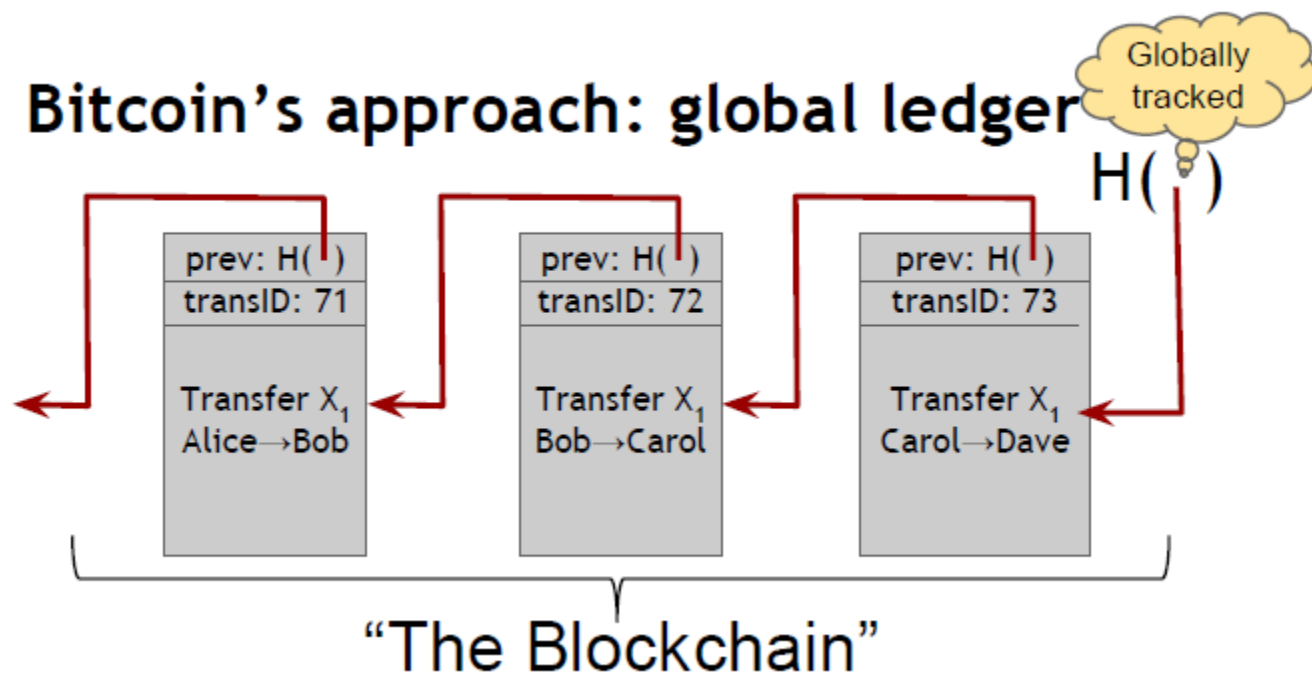


$H()$



optimization: put multiple transactions in the same block

Bitcoin's approach: global ledger




Scrooge publishes ledger of all transactions
(a blockchain, signed by Scrooge)



signed by pk_{Scrooge}

$H()$



Merkle tree of transactions
in each block

CreateCoins transaction creates new coins

transID: 73 type:CreateCoins		
coins created		
<i>num</i>	<i>value</i>	<i>recipient</i>
0	3.2	0x...
1	1.4	0x...
2	7.1	0x...

← coinID 73(0)

← coinID 73(1)

← coinID 73(2)

Valid, because I said so.



PayCoins transaction consumes (and destroys) some coins,
and creates new coins of the same total value

transID: 73 type:PayCoins		
consumed coinIDs: 68(1), 42(0), 72(3)		
coins created		
<i>num</i>	<i>value</i>	<i>recipient</i>
0	3.2	0x...
1	1.4	0x...
2	7.1	0x...
signatures		

Valid if:

- consumed coins valid,
- not already consumed,
- total value out = total value in, and
- signed by owners of all consumed coins

Immutable coins

Coins can't be transferred, subdivided, or combined.


But: you can get the same effect by using transactions
to subdivide: create new trans
consume your coin
pay out two new coins to yourself

Scrooge publishes ledger of all transactions
(a blockchain, signed by Scrooge)



signed by pk_{Scrooge}

$H()$

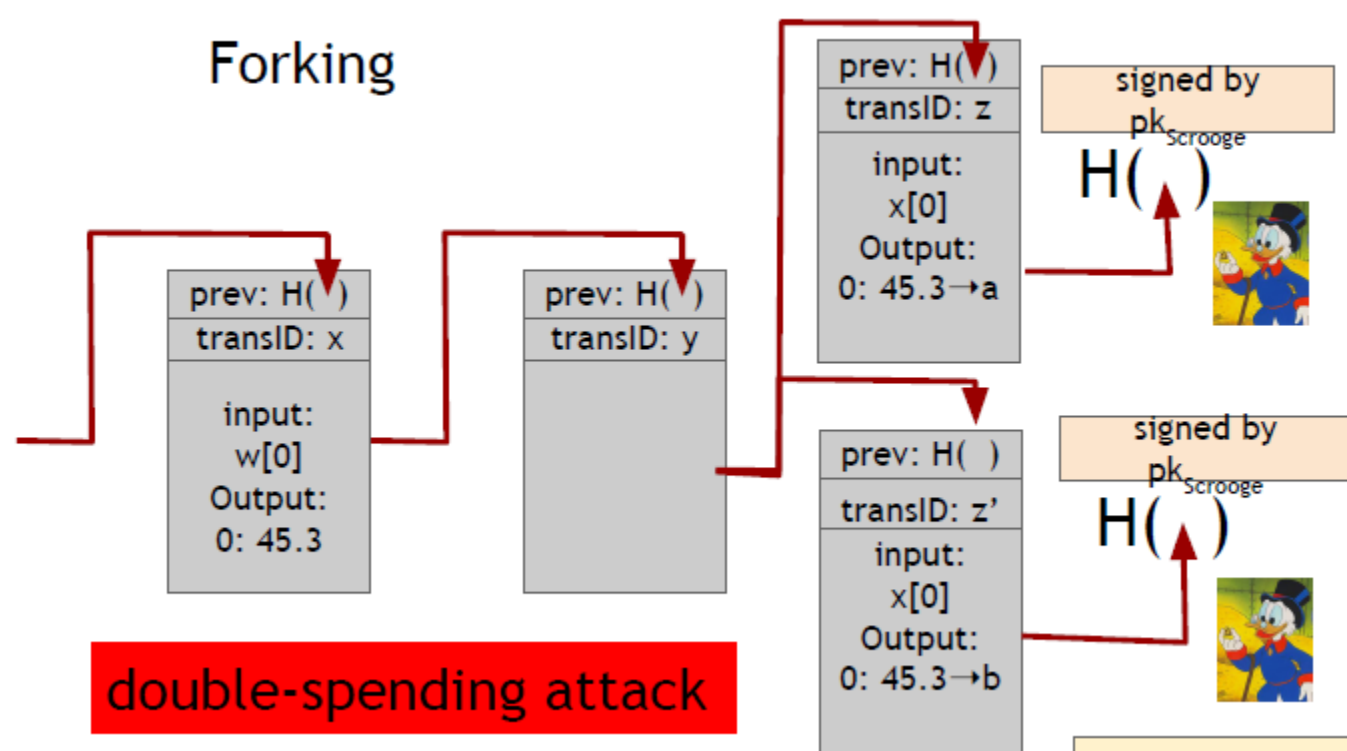


Merkle tree of transactions
in each block

Don't worry, I'm honest.



Forking



What if Scrooge is malicious?

Don't worry, I'm honest.



Crucial question:

Can we descroogify the
currency, and operate without
any central, trusted party?

The path to decentralization

- technology & incentive design



Who maintains the **ledger of transactions**? (and how?)

Who determines the **validity of transactions** to be included in the ledger?




Who creates **new Bitcoins**?

Mechanics of Bitcoin

What is blockchain, and why does it matter?



- A blockchain is a historical record of transactions, much like a database
- Blocks in a chain = pages in a book.
- Each page in a book contains:
 - The text: the story
 - Each page has information about itself: title of the book, chapter title, page number, etc. (e.g. the “metadata”)



“Bits on Blocks”, Blog by Antony Lewis, <https://bitsonblocks.net/>

- Similarly, in a blockchain, each block has:
 - A *header* which contains the data about the block: e.g. technical information, a reference to the previous block, and a digital fingerprint (aka “hash”) of the data contained in this block, among other things. This hash is important for ordering and block validation.
 - The *contents* of the block, e.g. information about the transaction(s)

Peer to Peer Network

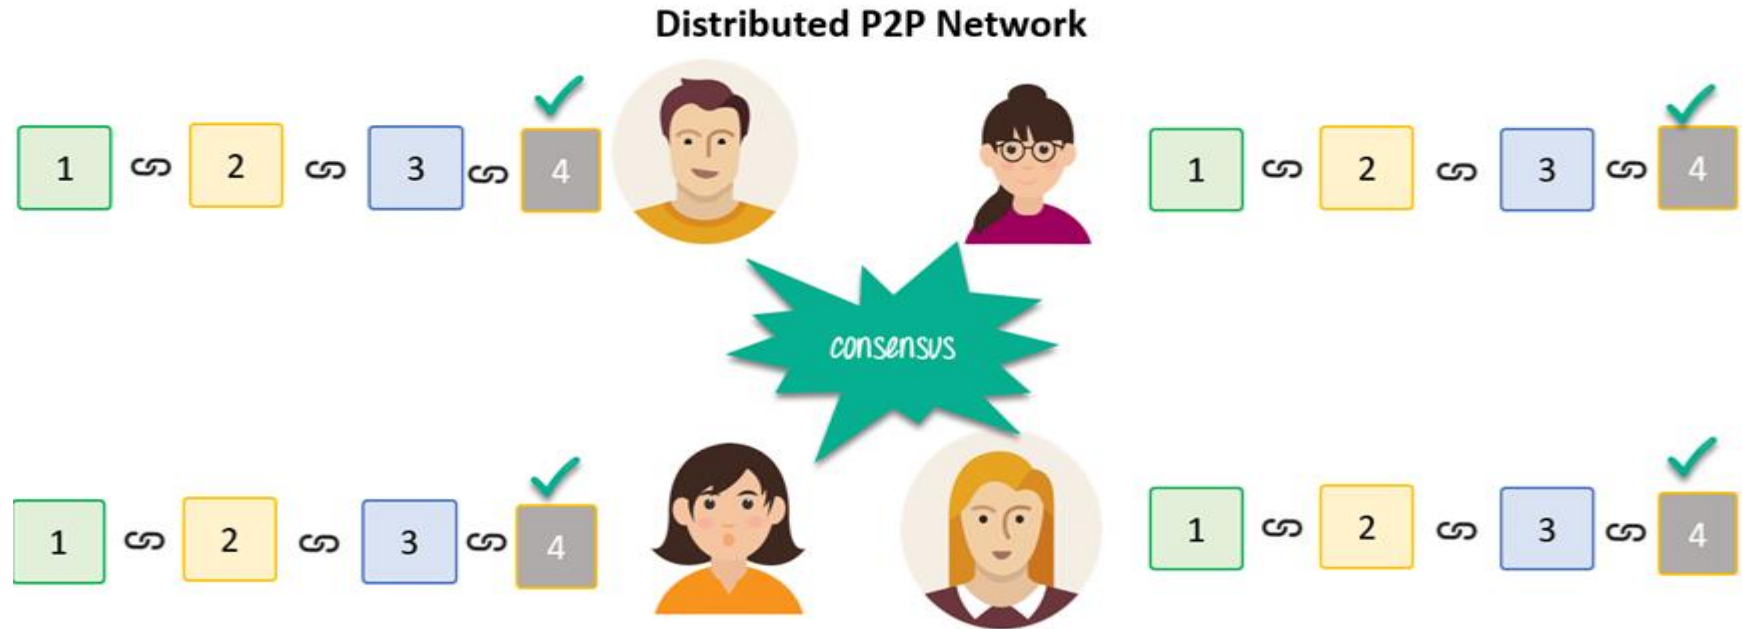


Peer to Peer Network

A distributed network architecture may be called a Peer-to-Peer (P-to-P, P2P,..) network, if the participants share a part of their own hardware resources (processing power, storage capacity, network link capacity, printers,..). These shared resources are necessary to provide the Service and content offered by the network (e.g. file sharing or shared workspaces for collaboration). They are accessible by other peers directly, without passing intermediary entities. The participants of such a network are thus resource (Service and content) providers as well as resource (Service and content) requestors (Servent-concept).

- Rüdiger Schollmeier, 2002

How to achieve consistency?



Public Keys as Identities

Why don't Bitcoin nodes have identities?

No Central Authority

Useful trick: public key == an identity

if you see sig such that $verify(pk, msg, sig) == true$,
think of it as

pk says, “[msg]”.

to “speak for” pk , you must know matching secret key sk

How to make a new identity

create a new, random key-pair (sk , pk)

pk is the public “name” you can use

[usually better to use $\text{Hash}(pk)$]

sk lets you “speak for” the identity

you control the identity, because only you know sk

if pk “looks random”, nobody needs to know who you are

Decentralized identity management

Public Key as Identity!

anybody can make a new identity at any time - make as many as you want!

These identities are called “addresses” in Bitcoin.

Pseudonymity is a goal of Bitcoin

Identity is hard in a P2P system – Sybil attack

Privacy

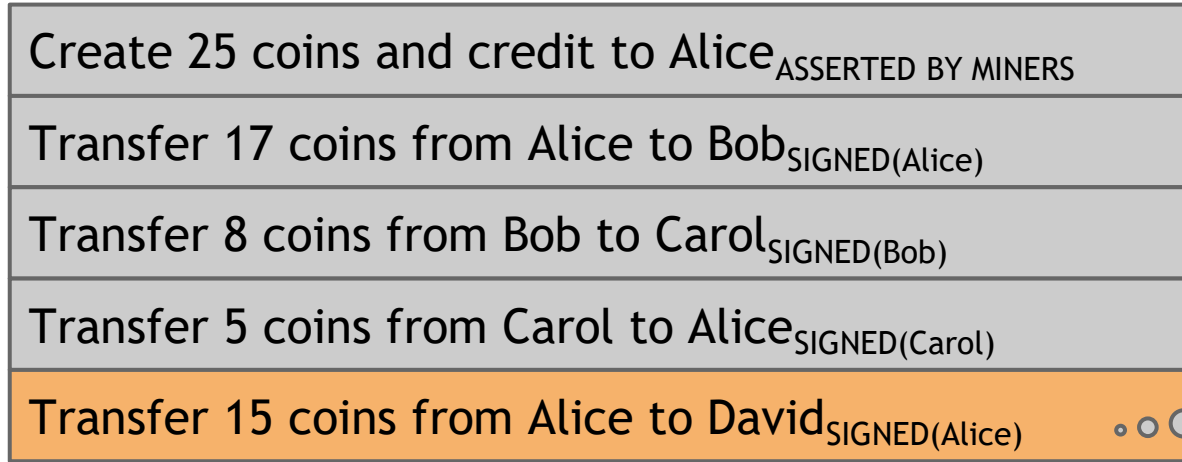
Addresses not directly connected to real-world identity.

But observer can link together an address's activity over time, make inferences.

Bitcoin transactions

An account-based ledger (*not* Bitcoin)

time

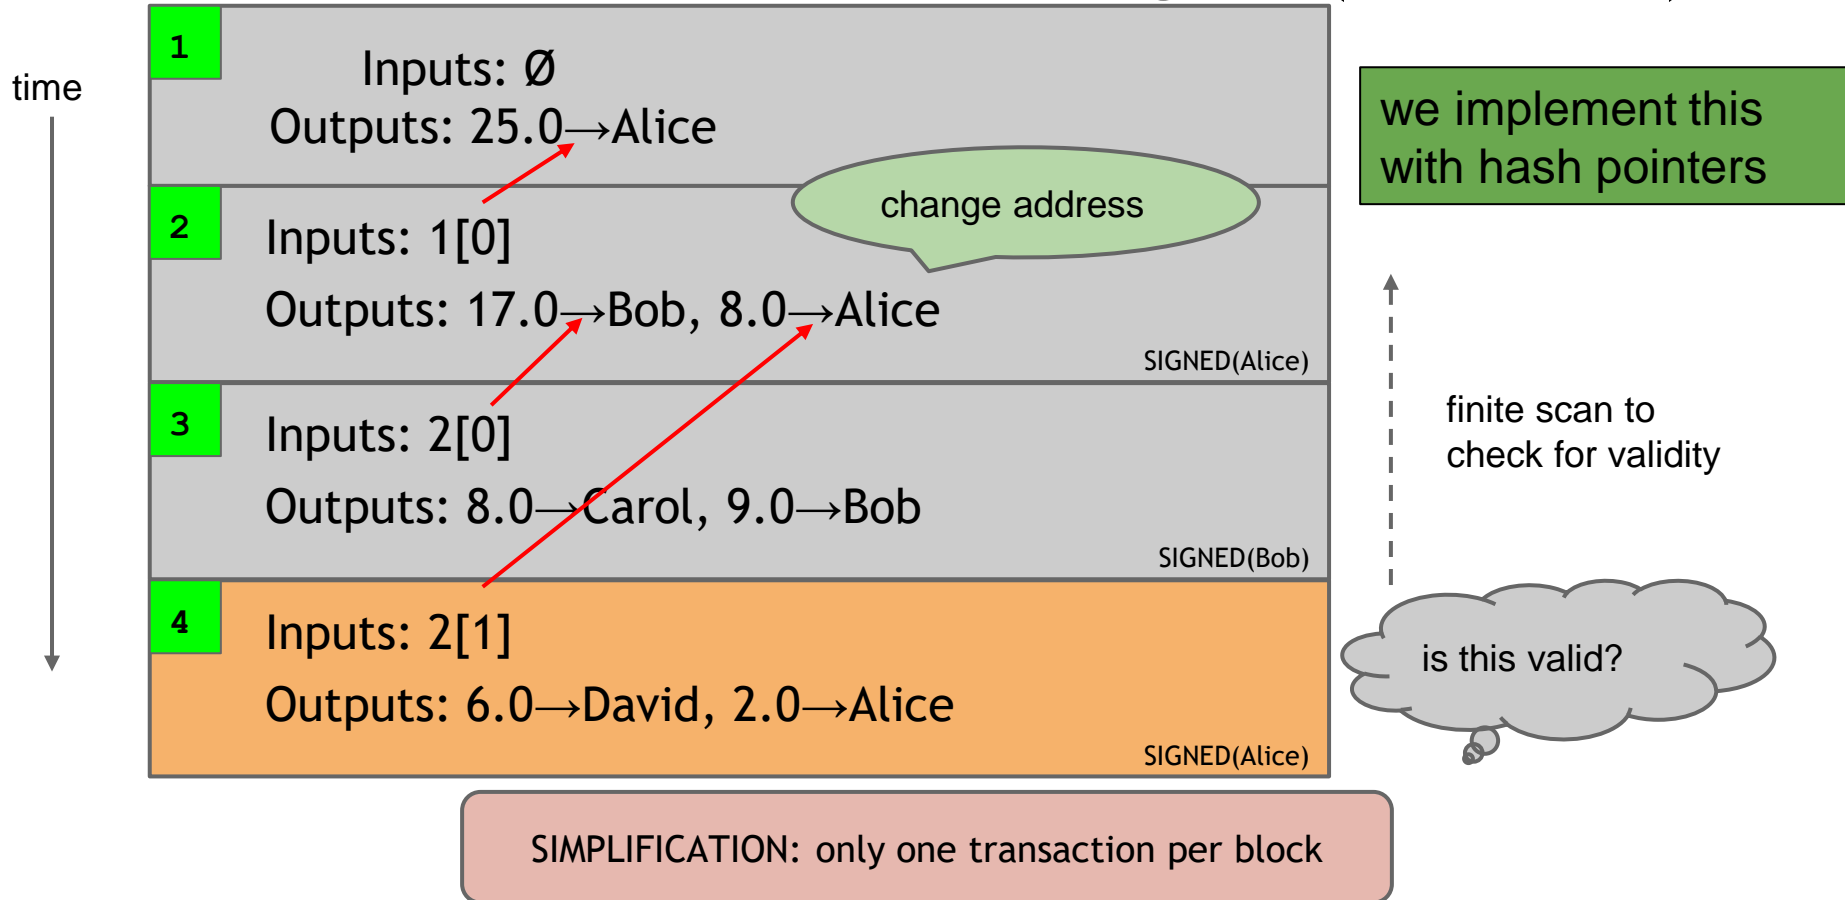


might need to
scan backwards
until genesis!

is this valid?

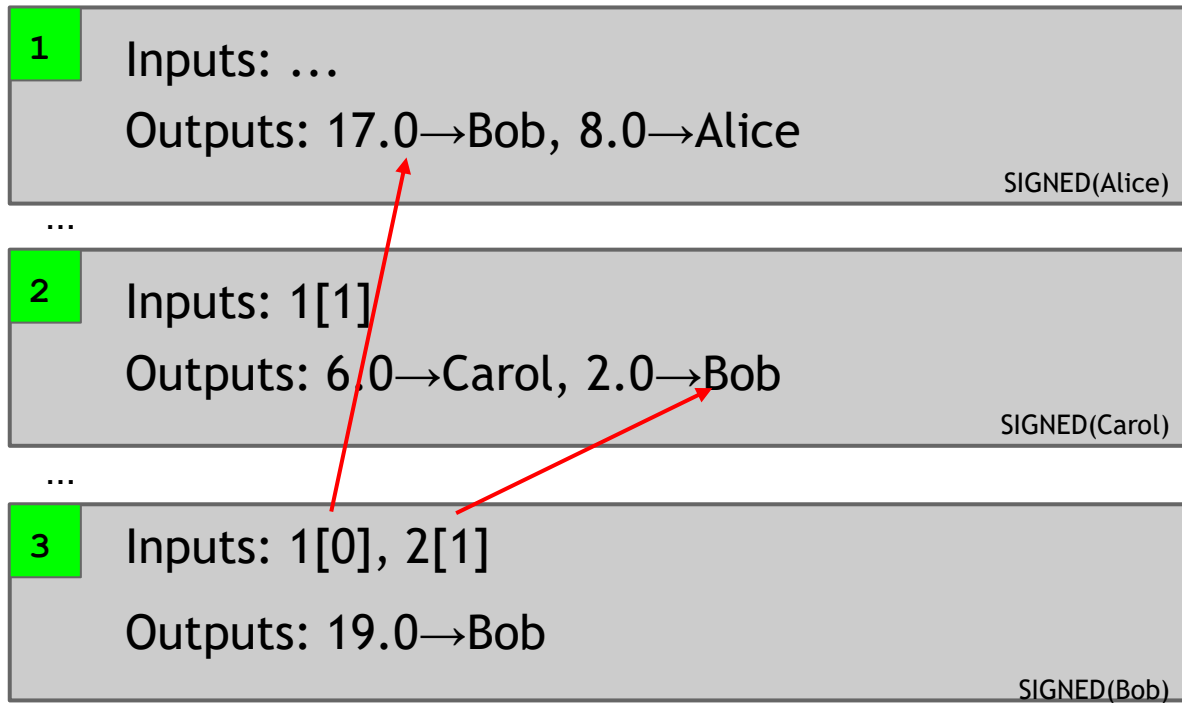
SIMPLIFICATION: only one transaction per block

A transaction-based ledger (Bitcoin)



Merging value

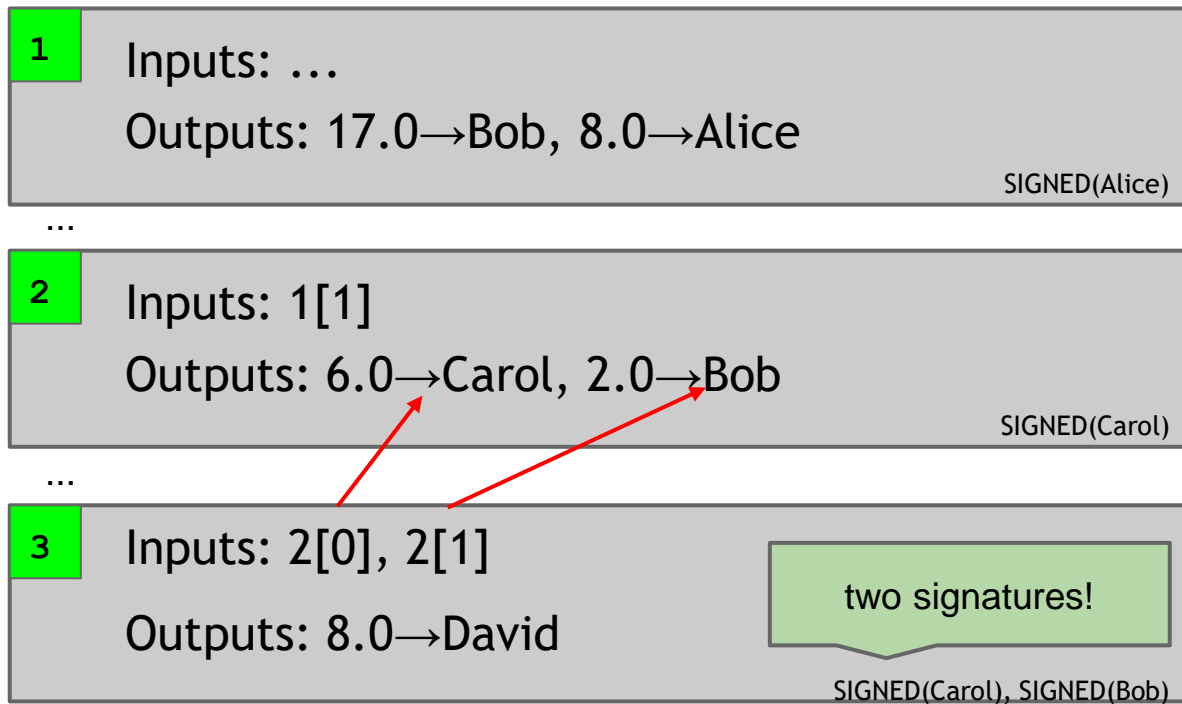
time



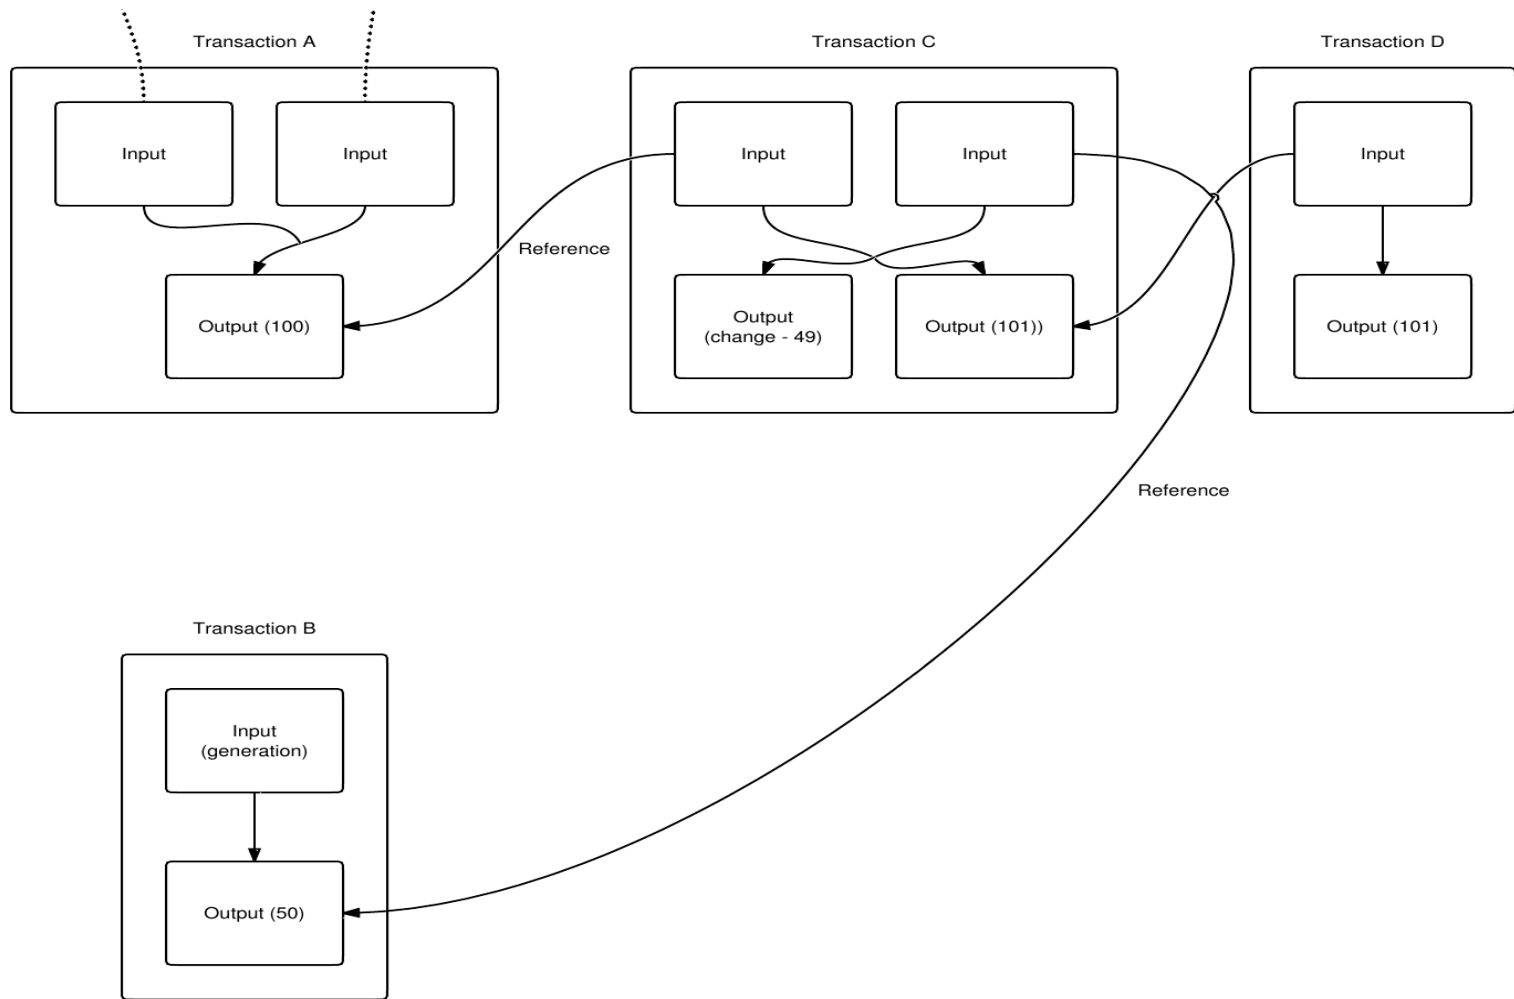
SIMPLIFICATION: only one transaction per block

Joint payments

time



SIMPLIFICATION: only one transaction per block



The real deal: a Bitcoin transaction

```
{
  "hash": "5a42590fbe0a90ee8e8747244d6c84f0db1a3a24e8f1b95b10c9e050990b8b6b",
  "ver": 1,
  "vin_sz": 2,
  "vout_sz": 1,
  "lock_time": 0,
  "size": 404,
  "in": [
    {
      "prev_out": {
        "hash": "3be4ac9728a0823cf5e2deb2e86fc0bd2aa503a91d307b42ba76117d79280260",
        "n": 0
      },
      "scriptSig": "30440..."
    },
    {
      "prev_out": {
        "hash": "7508e6ab259b4df0fd5147bab0c949d81473db4518f81afc5c3f52f91ff6b34e",
        "n": 0
      },
      "scriptSig": "3f3a4ce81...."
    }
  ],
  "out": [
    {
      "value": "10.12287097",
      "scriptPubKey": "OP_DUP OP_HASH160 69e02e18b5705a05dd6b28ed517716c894b3d42e OP_EQUALVERIFY OP_CHECKSIG"
    }
  ]
}
```

metadata

input(s)

output(s)

The real deal: transaction metadata

```
{
transaction hash { "hash":"5a42590...b8b6b",
housekeeping { "ver":1,
               "vin_sz":2,
               "vout_sz":1,
"not valid before" { "lock_time":0,
housekeeping { "size":404,
...
}
```

The real deal: transaction inputs

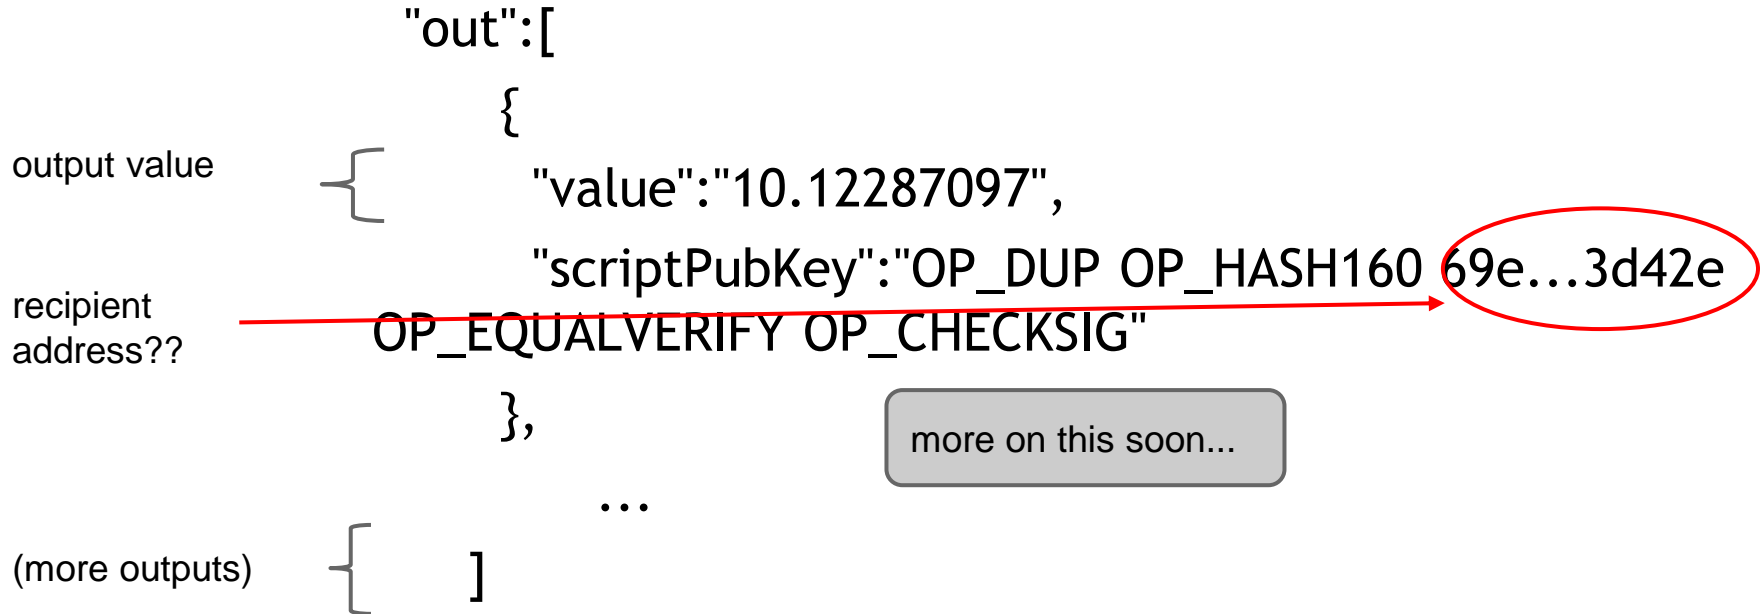
```
"in":[  
  {  
    "prev_out":{  
      "hash":"3be4...80260",  
      "n":0  
    },  
    "scriptSig":"30440....3f3a4ce81"  
  },  
  ...  
],
```

previous transaction {

signature {

(more inputs) {

The real deal: transaction outputs



Bitcoin scripts

Output “addresses” are really *scripts*

OP_DUP

OP_HASH160

69e02e18...

OP_EQUALVERIFY OP_CHECKSIG

Input “addresses” are *also* scripts

scriptSig

30440220...
0467d2c9...

scriptPubKey

OP_DUP
OP_HASH160
69e02e18...
OP_EQUALVERIFY OP_CHECKSIG


TO VERIFY: Concatenated script must execute completely with no errors

Bitcoin scripting language (“Script”)


Design goals

- Built for Bitcoin (inspired by Forth)
- Simple, compact
- Support for cryptography
- Stack-based
- Limits on time/memory
- No looping

I am not
impressed



Bitcoin script execution example



<sig> <pubKey> OP_DUP OP_HASH160 <pubKeyHash?> OP_EQUALVERIFY OP_CHECKSIG

OP_DUP	Duplicates the top item on the stack
OP_HASH160	Hashes twice: first using SHA-256 and then RIPEMD-160
OP_EQUALVERIFY	Returns true if the inputs are equal. Returns false and marks the transaction as invalid if they are unequal
OP_CHECKSIG	Checks that the input signature is a valid signature using the input public key for the hash of the current transaction
OP_CHECKMULTISIG	Checks that the k signatures on the transaction are valid signatures from k of the specified public keys.

Figure 3.6 a list of common Script instructions and their functionality.

Bitcoin script instructions

256 opcodes total (15 disabled, 75 reserved)

- Arithmetic
- If/then
- Logic/data handling
- Crypto!
 - Hashes
 - Signature verification
 - Multi-signature verification

OP_CHECKMULTISIG

- Built-in support for joint signatures
- Specify n public keys
- Specify t
- Verification requires t signatures



BUG ALERT: Extra data value
popped from the stack and ignored

Bitcoin scripts in practice (as of 2014)

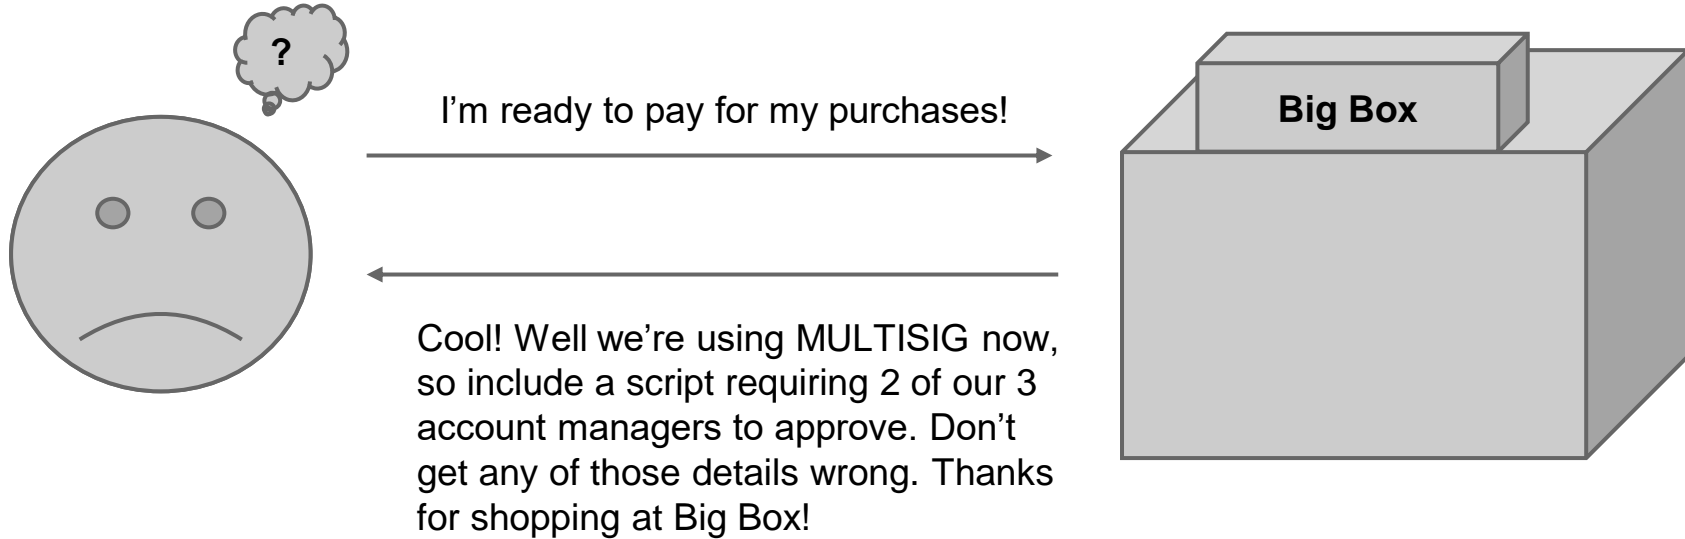
- Most nodes whitelist known scripts
- 99.9% are simple signature checks
- ~0.01% are MULTISIG
- ~0.01% are **Pay-to-Script-Hash**
- Remainder are errors, proof-of-burn

Proof-of-burn

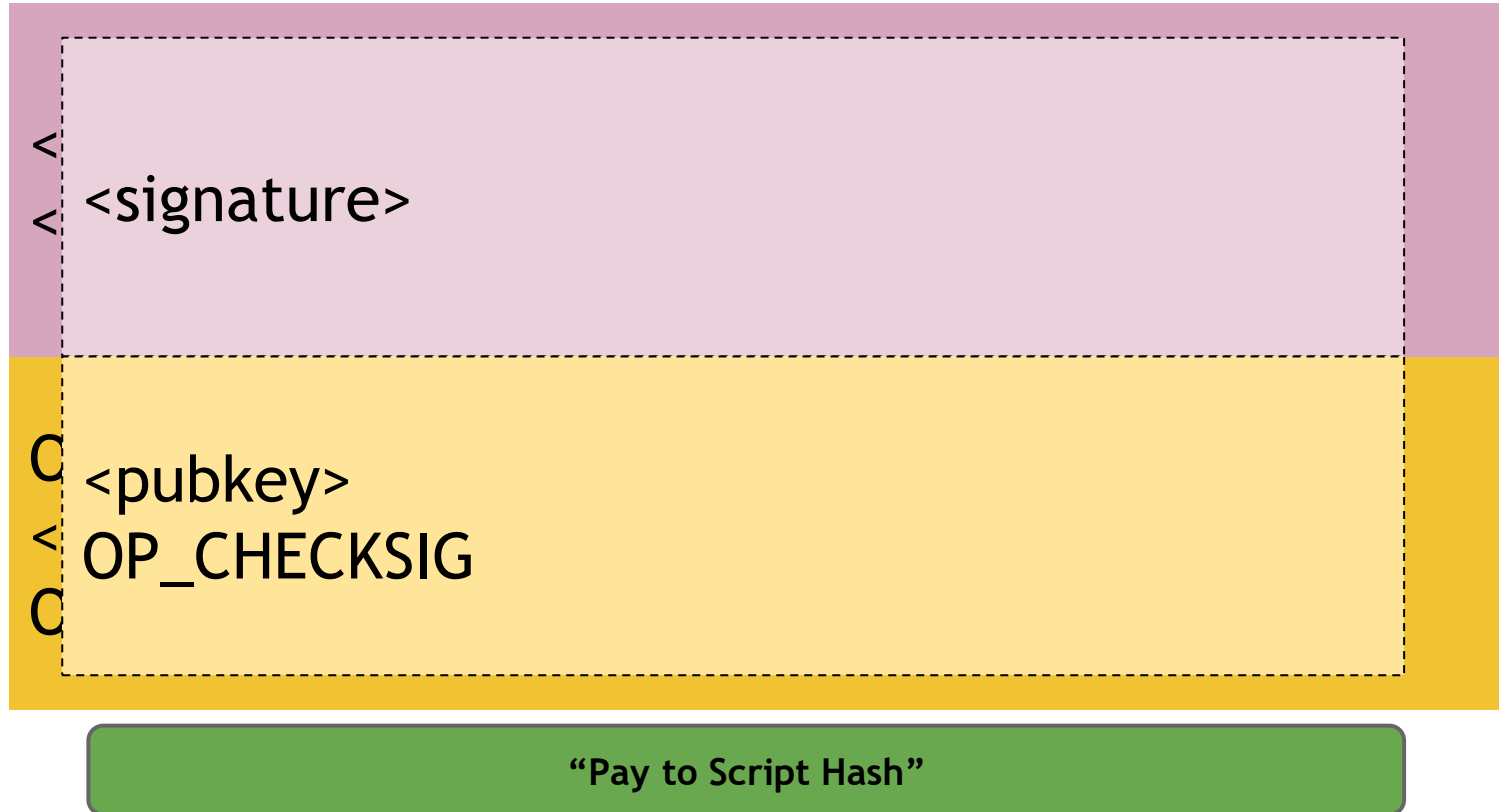
nothing's going to redeem that ☹️

OP_RETURN
<arbitrary data>

Should senders specify scripts?



Idea: use the hash of redemption script




Pay to script hash



I'm ready to pay for my purchases!

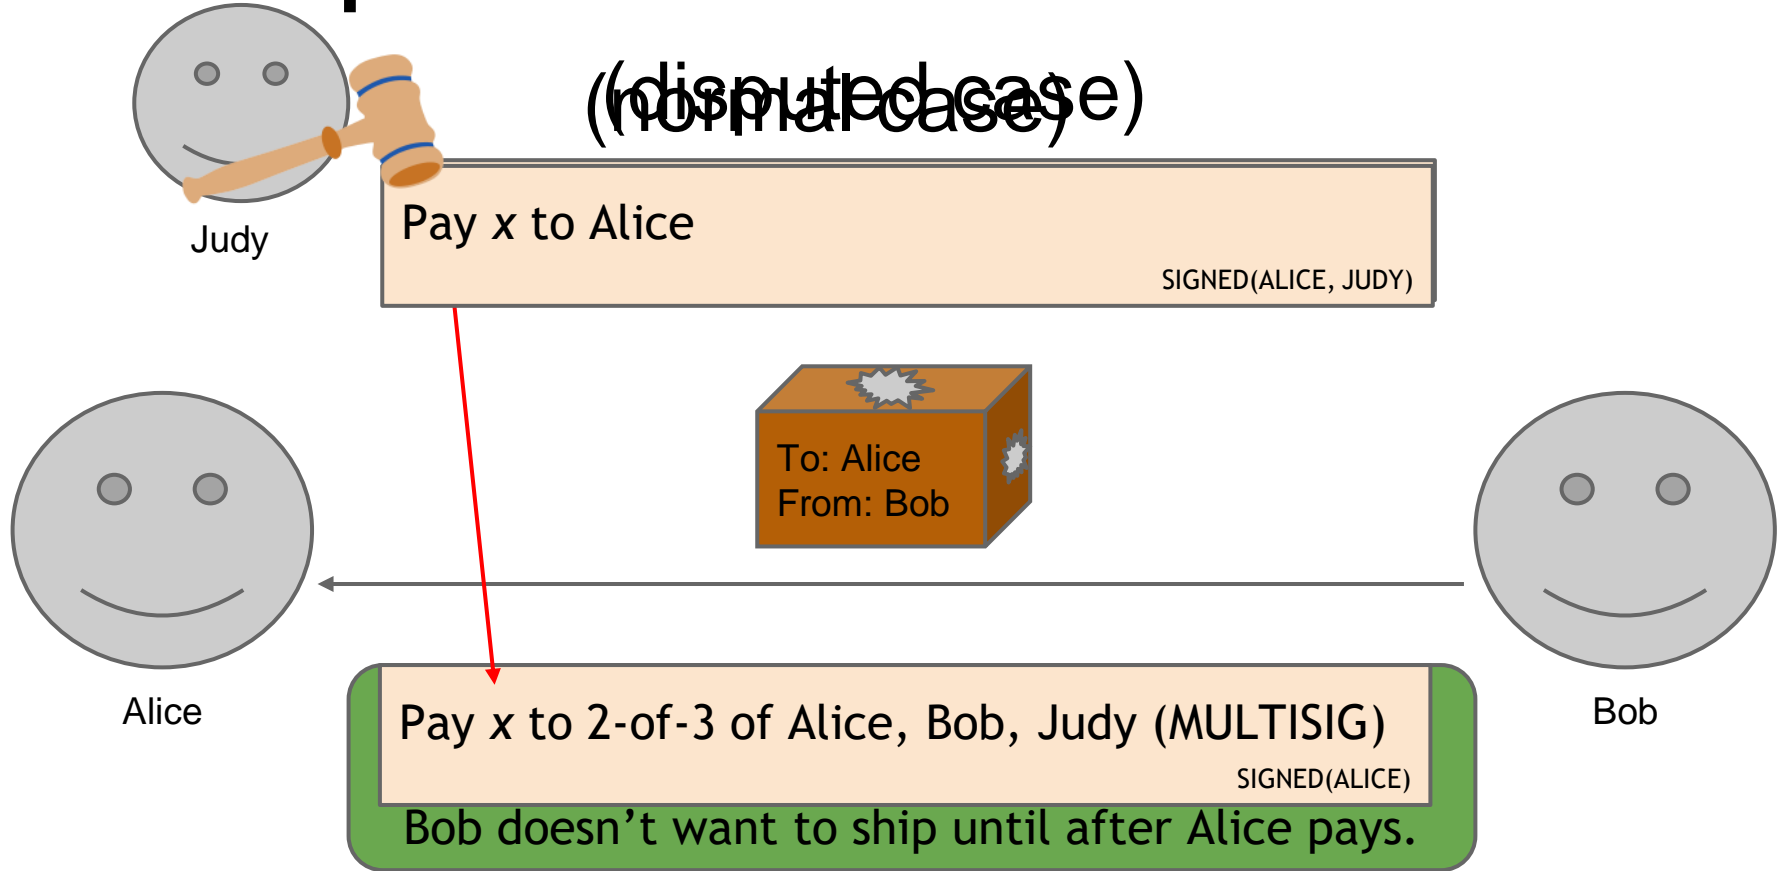


Great! Here's our address: 0x3454




Applications of Bitcoin scripts

Example 1: Escrow transactions



Example 2: Green addresses

InstaWallet, Mt. Gox
Collapsed!



Bank

You may Trust me. We do not do
double spend!

It's me, Alice! Could you make out
a green payment to Bob?



Alice

Pay x to Bob, y to Bank

No double spend

SIGNED(BANK)



Bob

PROBLEM: Alice wants to pay Bob.
Bob can't wait 6 verifications to guard against
double-spends, or is offline completely.

Example 3: Efficient micro-payments

What if Bob never signs??

all of these could
be double-
spends!

Input: x ; Pay 42 to Bob, 58 to Alice

SIGNED(ALICE) SIGNED(BOB)

...

Alice demands a timed refund transaction before starting

Input: x ; Pay 100 to Alice, LOCK until time t

SIGNED(ALICE) SIGNED(BOB)

I'm done!

Input: x ; Pay 03 to Bob, 97 to Alice

SIGNED(ALICE) _____

I'll publish!

Input: x ; Pay 02 to Bob, 98 to Alice

SIGNED(ALICE) _____

Input: x ; Pay 01 to Bob, 99 to Alice

SIGNED(ALICE) _____

PROBLEM: Alice wants to pay Bob for each

Input: y ; Pay 100 to Bob/Alice (MULTISIG)

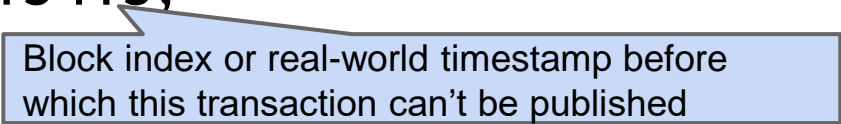
SIGNED(ALICE)

Alice

Bob

lock_time

```
{  
  "hash":"5a42590...b8b6b",  
  "ver":1,  
  "vin_sz":2,  
  "vout_sz":1,  
  "lock_time":315415,  
  "size":404,  
  ...  
}
```



Block index or real-world timestamp before
which this transaction can't be published

More advanced scripts

- Multiplayer lotteries
- Hash pre-image challenges
- Coin-swapping protocols

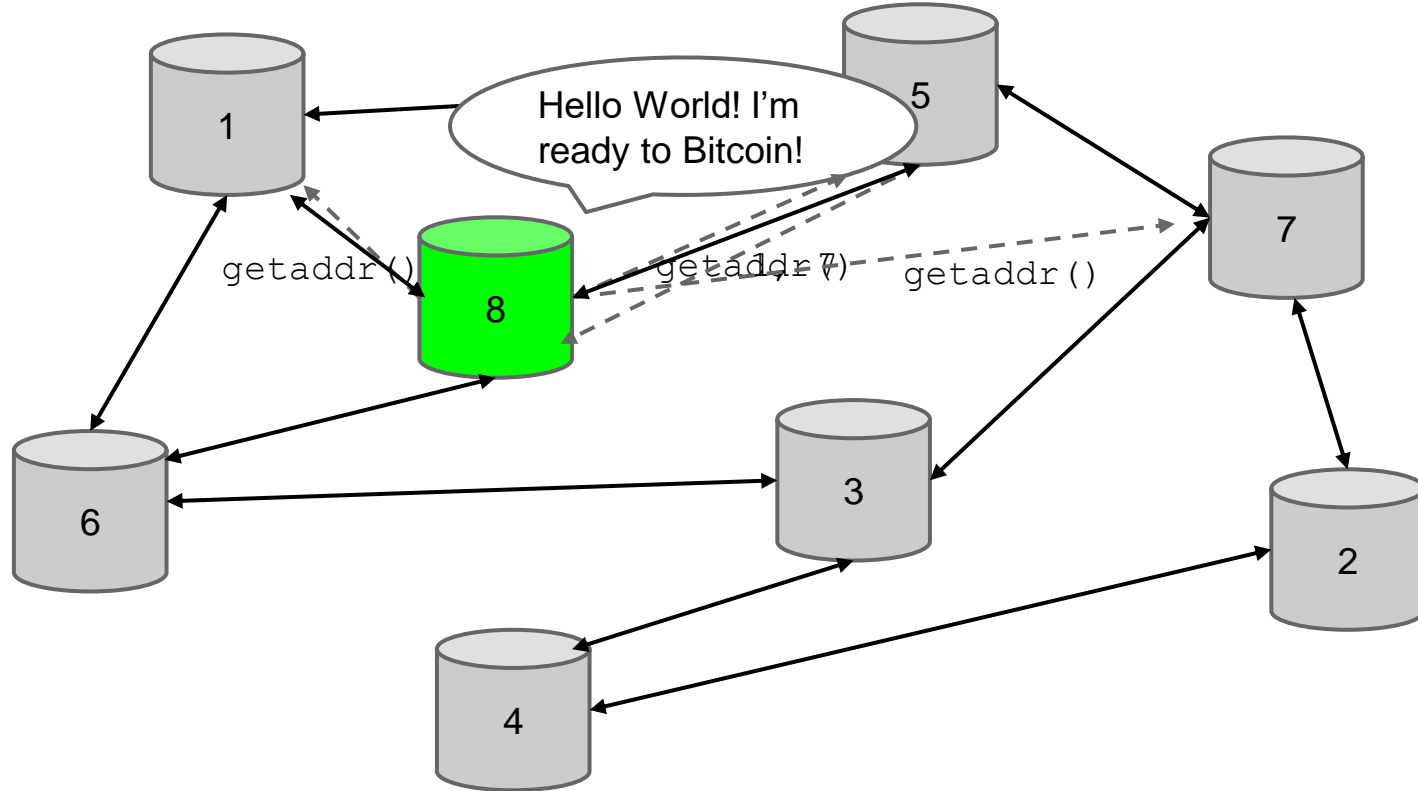
“Smart contracts”

The Bitcoin network


Bitcoin P2P network

- Ad-hoc protocol (runs on TCP port 8333)
- Ad-hoc network with random topology
- All nodes are equal
- New nodes can join at any time
 - Network Changes over time - dynamic
- No explicit way to leave network
 - Forget non-responding nodes after 3 hr

Joining the Bitcoin P2P network

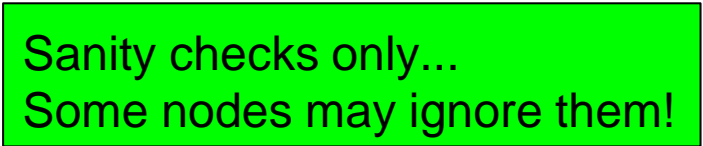


Transaction propagation (flooding)




Should I relay a proposed transaction?

- Transaction valid with current block chain
- (default) script matches a whitelist
 - Avoid unusual scripts
- Haven't seen before
 - Avoid infinite loops
- Doesn't conflict with others I've relayed
 - Avoid double-spends



Sanity checks only...
Some nodes may ignore them!

Nodes may differ on transaction pool



Race conditions

Transactions or blocks may *conflict*

- Default behavior: accept what you hear first
- Network position matters
- Miners may implement other logic!